

```
1 from execute_util import text, link, image
2 from facts import a100_flop_per_sec, h100_flop_per_sec
3 import torch.nn.functional as F
4 import timeit
5 import torch
6 from typing import Iterable
7 from torch import nn
8 import numpy as np
9 from lecture_util import article_link
10 from jaxtyping import Float
11 from einops import rearrange, einsum, reduce
12 from references import zero_2019
13
14
15 def main():
16     Last lecture: overview, tokenization
17
18     Overview of this lecture:
19     • We will discuss all the primitives needed to train a model.
20     • We will go bottom-up from tensors to models to optimizers to the training loop.
21     • We will pay close attention to efficiency (use of resources).
22
23     In particular, we will account for two types of resources:
24     • Memory (GB)
25     • Compute (FLOPs)
26
27     motivating_questions()
28
29     We will not go over the Transformer.
30     There are excellent expositions:
31     Assignment 1 handout
32     Mathematical description
33     Illustrated Transformer
34     Illustrated GPT-2
35     Instead, we'll work with simpler models.
36
37     What knowledge to take away:
38     • Mechanics: straightforward (just PyTorch)
39     • Mindset: resource accounting (remember to do it)
40     • Intuitions: broad strokes (no large models)
41
42     Memory accounting
43     tensors_basics()
44     tensors_memory()
45
46     Compute accounting
47     tensors_on_gpus()
48     tensor_operations()
49     tensor_einops()
50     tensor_operations_flops()
51     gradients_basics()
52     gradients_flops()
53
```

Models

```
module_parameters()
custom_model()
```

Training loop and best practices

```
note_about_randomness()
data_loading()
```

```
optimizer()
```

```
train_loop()
```

```
checkpointing()
```

```
mixed_precision_training()
```

```
def motivating_questions():
```

Let's do some napkin math.

Question: How long would it take to train a 70B parameter model on 15T tokens on 1024 H100s?

```
total_flops = 6 * 70e9 * 15e12 # @inspect total_flops
```

```
assert h100_flop_per_sec == 1979e12 / 2
```

```
mfu = 0.5
```

```
flops_per_day = h100_flop_per_sec * mfu * 1024 * 60 * 60 * 24 # @inspect flops_per_day
```

```
days = total_flops / flops_per_day # @inspect days
```

Question: What's the largest model that you can train on 8 H100s using AdamW (naively)?

```
h100_bytes = 80e9 # @inspect h100_bytes
```

```
bytes_per_parameter = 4 + 4 + (4 + 4) # parameters, gradients, optimizer state @inspect bytes_per_parameter
```

```
num_parameters = (h100_bytes * 8) / bytes_per_parameter # @inspect num_parameters
```

Caveat 1: we are naively using float32 for parameters and gradients. We could also use bf16 for parameters and gradients (2 + 2) and keep an extra float32 copy of the parameters (4). This doesn't save memory, but is faster. [\[Rajbhandari+ 2019\]](#)

Caveat 2: activations are not accounted for (depends on batch size and sequence length).

This is a rough back-of-the-envelope calculation.

```
def tensors_basics():
```

Tensors are the basic building block for storing everything: parameters, gradients, optimizer state, data, activations.

[\[PyTorch docs on tensors\]](#)

You can create tensors in multiple ways:

```
x = torch.tensor([[1., 2, 3], [4, 5, 6]]) # @inspect x
```

```
x = torch.zeros(4, 8) # 4x8 matrix of all zeros @inspect x
```

```
x = torch.ones(4, 8) # 4x8 matrix of all ones @inspect x
```

```
x = torch.randn(4, 8) # 4x8 matrix of iid Normal(0, 1) samples @inspect x
```

Allocate but don't initialize the values:

```
x = torch.empty(4, 8) # 4x8 matrix of uninitialized values @inspect x
```

...because you want to use some custom logic to set the values later

```
nn.init.trunc_normal_(x, mean=0, std=1, a=-2, b=2) # @inspect x
```

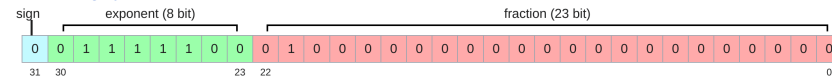
```
def tensors_memory():
```

Almost everything (parameters, gradients, activations, optimizer states) are stored as floating point numbers.

float32

[Wikipedia]

IEEE 754 single-precision 32-bit float



The float32 data type (also known as fp32 or single precision) is the default.

Traditionally, in scientific computing, float32 is the baseline; you could use double precision (float64) in some cases.

In deep learning, you can be a lot sloppier.

Let's examine memory usage of these tensors.

Memory is determined by the (i) number of values and (ii) data type of each value.

```
x = torch.zeros(4, 8) # @inspect x
assert x.dtype == torch.float32 # Default type
assert x.numel() == 4 * 8
assert x.element_size() == 4 # Float is 4 bytes
assert get_memory_usage(x) == 4 * 8 * 4 # 128 bytes
```

One matrix in the feedforward layer of GPT-3:

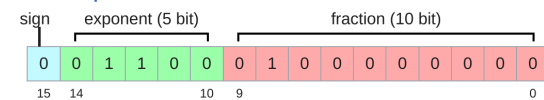
```
assert get_memory_usage(torch.empty(12288 * 4, 12288)) == 2304 * 1024 * 1024 # 2.3 GB
```

...which is a lot!

float16

[Wikipedia]

IEEE half-precision 16-bit float



The float16 data type (also known as fp16 or half precision) cuts down the memory.

```
x = torch.zeros(4, 8, dtype=torch.float16) # @inspect x
```

```
assert x.element_size() == 2
```

However, the dynamic range (especially for small numbers) isn't great.

```
x = torch.tensor([1e-8], dtype=torch.float16) # @inspect x
```

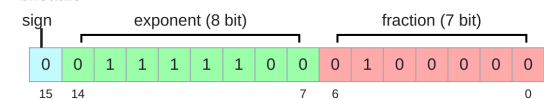
```
assert x == 0 # Underflow!
```

If this happens when you train, you can get instability.

bfloat16

[Wikipedia]

bfloat16



Google Brain developed bfloat (brain floating point) in 2018 to address this issue.

bfloat16 uses the same memory as float16 but has the same dynamic range as float32!

The only catch is that the resolution is worse, but this matters less for deep learning.

```
x = torch.tensor([1e-8], dtype=torch.bfloat16) # @inspect x
```

```
assert x != 0 # No underflow!
```

Let's compare the dynamic ranges and memory usage of the different data types:

```
float32_info = torch.finfo(torch.float32) # @inspect float32_info
```

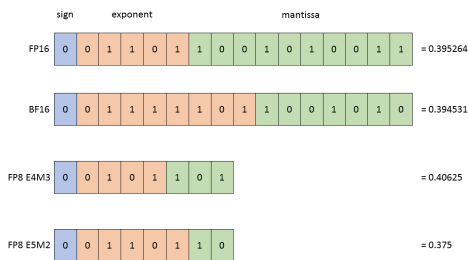
```
float16_info = torch.finfo(torch.float16) # @inspect float16_info
```

```
bfloat16_info = torch.finfo(torch.bfloat16) # @inspect bfloat16_info
```

fp8

In 2022, FP8 was standardized, motivated by machine learning workloads.

https://docs.nvidia.com/deeplearning/transformer-engine/user-guide/examples/fp8_primer.html



H100s support two variants of FP8: E4M3 (range [-448, 448]) and E5M2 ([-57344, 57344]).

Reference: [\[Micikevicius+ 2022\]](#)

Implications on training:

- Training with float32 works, but requires lots of memory.
- Training with fp8, float16 and even bfloat16 is risky, and you can get instability.
- Solution (later): use mixed precision training, see [mixed_precision_training](#)

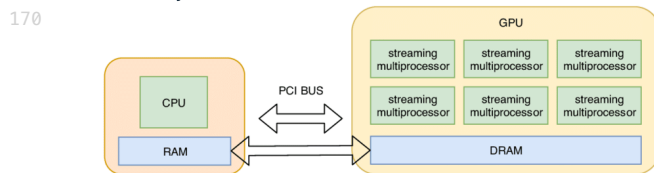
```
def tensors_on_gpus():
```

By default, tensors are stored in CPU memory.

```
x = torch.zeros(32, 32)
```

```
assert x.device == torch.device("cpu")
```

However, in order to take advantage of the massive parallelism of GPUs, we need to move them to GPU memory.



170

Let's first see if we have any GPUs.

```
if not torch.cuda.is_available():
```

```
    return
```

```
num_gpus = torch.cuda.device_count() # @inspect num_gpus
```

```
for i in range(num_gpus):
```

```
    properties = torch.cuda.get_device_properties(i) # @inspect properties
```

```
memory_allocated = torch.cuda.memory_allocated() # @inspect memory_allocated
```

```
text("Move the tensor to GPU memory (device 0).")
```

```
y = x.to("cuda:0")
```

```
assert y.device == torch.device("cuda", 0)
```

```
text("Or create a tensor directly on the GPU:")
```

```
z = torch.zeros(32, 32, device="cuda:0")
```

```
new_memory_allocated = torch.cuda.memory_allocated() # @inspect new_memory_allocated
```

```
memory_used = new_memory_allocated - memory_allocated # @inspect memory_used
```

```
assert memory_used == 2 * (32 * 32 * 4) # 2 32x32 matrices of 4-byte floats
```

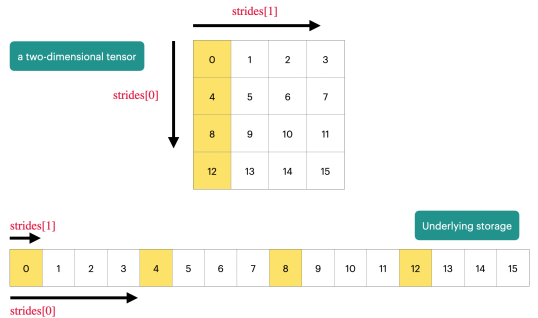
```
def tensor_operations():
```

Most tensors are created from performing operations on other tensors.

Each operation has some memory and compute consequence.

```

198
199 tensor_storage()
200 tensor_slicing()
201 tensor_elementwise()
202 tensor_matmul()
203
204
205 def tensor_storage():
206     What are tensors in PyTorch?
207     PyTorch tensors are pointers into allocated memory
208     ...with metadata describing how to get to any element of the tensor.
209 
```



```

210 [PyTorch docs]
211 x = torch.tensor([
212     [0., 1, 2, 3],
213     [4, 5, 6, 7],
214     [8, 9, 10, 11],
215     [12, 13, 14, 15],
216 ])
217
218 To go to the next row (dim 0), skip 4 elements in storage.
219 assert x.stride(0) == 4
220
221 To go to the next column (dim 1), skip 1 element in storage.
222 assert x.stride(1) == 1
223
224 To find an element:
225 r, c = 1, 2
226 index = r * x.stride(0) + c * x.stride(1) # @inspect index
227 assert index == 6
228
229 
```

```

230 def tensor_slicing():
231     x = torch.tensor([[1., 2, 3], [4, 5, 6]]) # @inspect x
232 
```

Many operations simply provide a different **view** of the tensor.
This does not make a copy, and therefore mutations in one tensor affects the other.

```

236 Get row 0:
237 y = x[0] # @inspect y
238 assert torch.equal(y, torch.tensor([1., 2, 3]))
239 assert same_storage(x, y)
240 
```

```

241 Get column 1:
242 y = x[:, 1] # @inspect y
243 assert torch.equal(y, torch.tensor([2, 5]))
244 assert same_storage(x, y)
245 
```

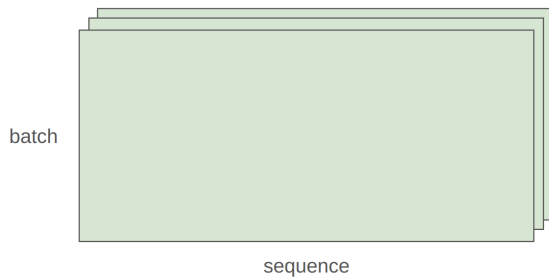
View 2x3 matrix as 3x2 matrix:

```

247 y = x.view(3, 2) # @inspect y
248 assert torch.equal(y, torch.tensor([[1, 2], [3, 4], [5, 6]]))
249 assert same_storage(x, y)
250
251 Transpose the matrix:
252 y = x.transpose(1, 0) # @inspect y
253 assert torch.equal(y, torch.tensor([[1, 4], [2, 5], [3, 6]]))
254 assert same_storage(x, y)
255
256 Check that mutating x also mutates y.
257 x[0][0] = 100 # @inspect x, @inspect y
258 assert y[0][0] == 100
259
260 Note that some views are non-contiguous entries, which means that further views aren't possible.
261 x = torch.tensor([[1., 2, 3], [4, 5, 6]]) # @inspect x
262 y = x.transpose(1, 0) # @inspect y
263 assert not y.is_contiguous()
264 try:
265     y.view(2, 3)
266     assert False
267 except RuntimeError as e:
268     assert "view size is not compatible with input tensor's size and stride" in str(e)
269
270 One can enforce a tensor to be contiguous first:
271 y = x.transpose(1, 0).contiguous().view(2, 3) # @inspect y
272 assert not same_storage(x, y)
273 Views are free, copying take both (additional) memory and compute.
274
275
276 def tensor_elementwise():
277     These operations apply some operation to each element of the tensor
278     ...and return a (new) tensor of the same shape.
279
280     x = torch.tensor([1, 4, 9])
281     assert torch.equal(x.pow(2), torch.tensor([1, 16, 81]))
282     assert torch.equal(x.sqrt(), torch.tensor([1, 2, 3]))
283     assert torch.equal(x.rsqrt(), torch.tensor([1, 1 / 2, 1 / 3])) # i -> 1/sqrt(x_i)
284
285     assert torch.equal(x + x, torch.tensor([2, 8, 18]))
286     assert torch.equal(x * 2, torch.tensor([2, 8, 18]))
287     assert torch.equal(x / 0.5, torch.tensor([2, 8, 18]))
288
289     triu takes the upper triangular part of a matrix.
290     x = torch.ones(3, 3).triu() # @inspect x
291     assert torch.equal(x, torch.tensor([
292         [1, 1, 1],
293         [0, 1, 1],
294         [0, 0, 1]],
295     ))
296     This is useful for computing an causal attention mask, where  $M[i, j]$  is the contribution of  $i$  to  $j$ .
297
298
299 def tensor_matmul():
300     Finally, the bread and butter of deep learning: matrix multiplication.
301     x = torch.ones(16, 32)
302     w = torch.ones(32, 2)
303     y = x @ w
304     assert y.size() == torch.Size([16, 2])

```

In general, we perform operations for every example in a batch and token in a sequence.



```
x = torch.ones(4, 8, 16, 32)
```

```
w = torch.ones(32, 2)
```

```
y = x @ w
```

```
assert y.size() == torch.Size([4, 8, 16, 2])
```

In this case, we iterate over values of the first 2 dimensions of x and multiply by w.

```
def tensor_einops():
```

```
    einops_motivation()
```

Einops is a library for manipulating tensors where dimensions are named.

It is inspired by Einstein summation notation (Einstein, 1916).

[\[Einops tutorial\]](#)

```
jaxtyping_basics()
```

```
einops_einsum()
```

```
einops_reduce()
```

```
einops_rearrange()
```

```
def einops_motivation():
```

Traditional PyTorch code:

```
x = torch.ones(2, 2, 3) # batch, sequence, hidden @inspect x
```

```
y = torch.ones(2, 2, 3) # batch, sequence, hidden @inspect y
```

```
z = x @ y.transpose(-2, -1) # batch, sequence, sequence @inspect z
```

Easy to mess up the dimensions (what is -2, -1?)...

```
def jaxtyping_basics():
```

How do you keep track of tensor dimensions?

Old way:

```
x = torch.ones(2, 2, 1, 3) # batch seq heads hidden @inspect x
```

New (jaxtyping) way:

```
x: Float[torch.Tensor, "batch seq heads hidden"] = torch.ones(2, 2, 1, 3) # @inspect x
```

Note: this is just documentation (no enforcement).

```
def einops_einsum():
```

Einsum is generalized matrix multiplication with good bookkeeping.

Define two tensors:

```
x: Float[torch.Tensor, "batch seq1 hidden"] = torch.ones(2, 3, 4) # @inspect x
```

```
y: Float[torch.Tensor, "batch seq2 hidden"] = torch.ones(2, 3, 4) # @inspect y
```

Old way:

```

355     z = x @ y.transpose(-2, -1) # batch, sequence, sequence @inspect z
356
357     New (einops) way:
358     z = einsum(x, y, "batch seq1 hidden, batch seq2 hidden -> batch seq1 seq2") # @inspect z
359     Dimensions that are not named in the output are summed over.
360
361     Or can use ... to represent broadcasting over any number of dimensions:
362     z = einsum(x, y, "... seq1 hidden, ... seq2 hidden -> ... seq1 seq2") # @inspect z
363
364
365 def einops_reduce():
366     You can reduce a single tensor via some operation (e.g., sum, mean, max, min).
367     x: Float[torch.Tensor, "batch seq hidden"] = torch.ones(2, 3, 4) # @inspect x
368
369     Old way:
370     y = x.mean(dim=-1) # @inspect y
371
372     New (einops) way:
373     y = reduce(x, "... hidden -> ...", "sum") # @inspect y
374
375
376 def einops_rearrange():
377     Sometimes, a dimension represents two dimensions
378     ...and you want to operate on one of them.
379
380     x: Float[torch.Tensor, "batch seq total_hidden"] = torch.ones(2, 3, 8) # @inspect x
381     ...where total_hidden is a flattened representation of heads * hidden1
382     w: Float[torch.Tensor, "hidden1 hidden2"] = torch.ones(4, 4)
383
384     Break up total_hidden into two dimensions (heads and hidden1):
385     x = rearrange(x, "... (heads hidden1) -> ... heads hidden1", heads=2) # @inspect x
386
387     Perform the transformation by w:
388     x = einsum(x, w, "... hidden1, hidden1 hidden2 -> ... hidden2") # @inspect x
389
390     Combine heads and hidden2 back together:
391     x = rearrange(x, "... heads hidden2 -> ... (heads hidden2)") # @inspect x
392
393
394 def tensor_operations_flops():
395     Having gone through all the operations, let us examine their computational cost.
396
397     A floating-point operation (FLOP) is a basic operation like addition (x + y) or multiplication (x y).
398
399     Two terribly confusing acronyms (pronounced the same!):
400     • FLOPs: floating-point operations (measure of computation done)
401     • FLOP/s: floating-point operations per second (also written as FLOPS), which is used to measure the speed
      of hardware.
402
403
404 Intuitions
405     Training GPT-3 (2020) took 3.14e23 FLOPs. \[article\]
406     Training GPT-4 (2023) is speculated to take 2e25 FLOPs \[article\]
407     US executive order: any foundation model trained with >= 1e26 FLOPs must be reported to the government
      (revoked in 2025)
408
409     A100 has a peak performance of 312 teraFLOP/s \[spec\]
410     assert a100_flop_per_sec == 312e12

```


H100 has a peak performance of 1979 teraFLOP/s with sparsity, 50% without [\[spec\]](#)
`assert h100_flop_per_sec == 1979e12 / 2`

8 H100s for 2 weeks:

```
total_flops = 8 * (60 * 60 * 24 * 7) * h100_flop_per_sec # @inspect total_flops
```

Linear model

As motivation, suppose you have a linear model.

- We have n points
- Each point is d -dimensional
- The linear model maps each d -dimensional vector to a k outputs

```
if torch.cuda.is_available():
    B = 16384 # Number of points
    D = 32768 # Dimension
    K = 8192 # Number of outputs
else:
    B = 1024
    D = 256
    K = 64
```

```
device = get_device()
x = torch.ones(B, D, device=device)
w = torch.randn(D, K, device=device)
y = x @ w
```

We have one multiplication ($x[i][j] * w[j][k]$) and one addition per (i, j, k) triple.

```
actual_num_flops = 2 * B * D * K # @inspect actual_num_flops
```

FLOPs of other operations

- Elementwise operation on a $m \times n$ matrix requires $O(mn)$ FLOPs.
- Addition of two $m \times n$ matrices requires mn FLOPs.

In general, no other operation that you'd encounter in deep learning is as expensive as matrix multiplication for large enough matrices.

Interpretation:

- B is the number of data points
- $(D K)$ is the number of parameters
- FLOPs for forward pass is $2 (\# \text{ tokens}) (\# \text{ parameters})$

It turns out this generalizes to Transformers (to a first-order approximation).

How do our FLOPs calculations translate to wall-clock time (seconds)?

Let us time it!

```
actual_time = time_matmul(x, w) # @inspect actual_time
actual_flop_per_sec = actual_num_flops / actual_time # @inspect actual_flop_per_sec
```

Each GPU has a specification sheet that reports the peak performance.

- A100 [\[spec\]](#)
- H100 [\[spec\]](#)

Note that the FLOP/s depends heavily on the data type!

```
promised_flop_per_sec = get_promised_flop_per_sec(device, x.dtype) # @inspect promised_flop_per_sec
```

Model FLOPs utilization (MFU)

Definition: $(\text{actual FLOP/s}) / (\text{promised FLOP/s})$ [ignore communication/overhead]

```
mfu = actual_flop_per_sec / promised_flop_per_sec # @inspect mfu
```

Usually, MFU of ≥ 0.5 is quite good (and will be higher if matmuls dominate)

```

467 Let's do it with bfloat16:
468 x = x.to(torch.bfloat16)
469 w = w.to(torch.bfloat16)
470 bf16_actual_time = time_matmul(x, w) # @inspect bf16_actual_time
471 bf16_actual_flop_per_sec = actual_num_flops / bf16_actual_time # @inspect bf16_actual_flop_per_sec
472 bf16_promised_flop_per_sec = get_promised_flop_per_sec(device, x.dtype) # @inspect bf16_promised_flop_per_sec
473 bf16_mfu = bf16_actual_flop_per_sec / bf16_promised_flop_per_sec # @inspect bf16_mfu
474 Note: comparing bfloat16 to float32, the actual FLOP/s is higher.
475 The MFU here is rather low, probably because the promised FLOPs is a bit optimistic.
476
477

```

Summary

- Matrix multiplications dominate: (2 m n p) FLOPs
- FLOP/s depends on hardware (H100 >> A100) and data type (bfloat16 >> float32)
- Model FLOPs utilization (MFU): (actual FLOP/s) / (promised FLOP/s)

```

482
483 def gradients_basics():

```

So far, we've constructed tensors (which correspond to either parameters or data) and passed them through operations (forward).

Now, we're going to compute the gradient (backward).

As a simple example, let's consider the simple linear model:

$$y = 0.5 (x * w - 5)^2$$

Forward pass: compute loss

```

491 x = torch.tensor([1., 2, 3])
492 w = torch.tensor([1., 1, 1], requires_grad=True) # Want gradient
493 pred_y = x @ w
494 loss = 0.5 * (pred_y - 5).pow(2)
495

```

Backward pass: compute gradients

```

497 loss.backward()
498 assert loss.grad is None
499 assert pred_y.grad is None
500 assert x.grad is None
501 assert torch.equal(w.grad, torch.tensor([1, 2, 3]))
502
503

```

```

504 def gradients_flops():

```

Let us do count the FLOPs for computing gradients.

Revisit our linear model

```

508 if torch.cuda.is_available():
509     B = 16384 # Number of points
510     D = 32768 # Dimension
511     K = 8192 # Number of outputs
512 else:
513     B = 1024
514     D = 256
515     K = 64
516

```

```

517 device = get_device()

```

```

518 x = torch.ones(B, D, device=device)
519 w1 = torch.randn(D, D, device=device, requires_grad=True)
520 w2 = torch.randn(D, K, device=device, requires_grad=True)
521

```

Model: $x \rightarrow w_1 \rightarrow h_1 \rightarrow w_2 \rightarrow h_2 \rightarrow \text{loss}$

```

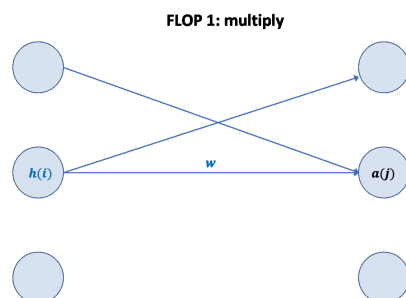
523 h1 = x @ w1

```

```

524 h2 = h1 @ w2
525 loss = h2.pow(2).mean()
526
527 Recall the number of forward FLOPs: tensor\_operations\_flops
528 • Multiply  $x[i][j] * w1[j][k]$ 
529 • Add to  $h1[i][k]$ 
530 • Multiply  $h1[i][j] * w2[j][k]$ 
531 • Add to  $h2[i][k]$ 
532 num_forward_flops = (2 * B * D * D) + (2 * B * D * K) # @inspect num_forward_flops
533
534 How many FLOPs is running the backward pass?
535 h1.retain_grad() # For debugging
536 h2.retain_grad() # For debugging
537 loss.backward()
538
539 Recall model:  $x \rightarrow w1 \rightarrow h1 \rightarrow w2 \rightarrow h2 \rightarrow \text{loss}$ 
540
541 •  $h1.\text{grad} = d \text{ loss} / d h1$ 
542 •  $h2.\text{grad} = d \text{ loss} / d h2$ 
543 •  $w1.\text{grad} = d \text{ loss} / d w1$ 
544 •  $w2.\text{grad} = d \text{ loss} / d w2$ 
545
546 Focus on the parameter  $w2$ .
547 Invoke the chain rule.
548
549 num_backward_flops = 0 # @inspect num_backward_flops
550
551  $w2.\text{grad}[j,k] = \sum_i h1[i,j] * h2.\text{grad}[i,k]$ 
552 assert  $w2.\text{grad}.\text{size}() == \text{torch.Size}([D, K])$ 
553 assert  $h1.\text{size}() == \text{torch.Size}([B, D])$ 
554 assert  $h2.\text{grad}.\text{size}() == \text{torch.Size}([B, K])$ 
555 For each  $(i, j, k)$ , multiply and add.
556 num_backward_flops += 2 * B * D * K # @inspect num_backward_flops
557
558  $h1.\text{grad}[i,j] = \sum_k w2[j,k] * h2.\text{grad}[i,k]$ 
559 assert  $h1.\text{grad}.\text{size}() == \text{torch.Size}([B, D])$ 
560 assert  $w2.\text{size}() == \text{torch.Size}([D, K])$ 
561 assert  $h2.\text{grad}.\text{size}() == \text{torch.Size}([B, K])$ 
562 For each  $(i, j, k)$ , multiply and add.
563 num_backward_flops += 2 * B * D * K # @inspect num_backward_flops
564
565 This was for just  $w2$  ( $D*K$  parameters).
566 Can do it for  $w1$  ( $D*D$  parameters) as well (though don't need  $x.\text{grad}$ ).
567 num_backward_flops += (2 + 2) * B * D * D # @inspect num_backward_flops
568
569 A nice graphical visualization: \[article\]
570

```



Putting it together:

- Forward pass: 2 (# data points) (# parameters) FLOPs
- Backward pass: 4 (# data points) (# parameters) FLOPs
- Total: 6 (# data points) (# parameters) FLOPs

```
def module_parameters():
```

```
    input_dim = 16384
```

```
    output_dim = 32
```

Model parameters are stored in PyTorch as `nn.Parameter` objects.

```
w = nn.Parameter(torch.randn(input_dim, output_dim))
```

```
assert isinstance(w, torch.Tensor) # Behaves like a tensor
```

```
assert type(w.data) == torch.Tensor # Access the underlying tensor
```

Parameter initialization

Let's see what happens.

```
x = nn.Parameter(torch.randn(input_dim))
```

```
output = x @ w # @inspect output
```

```
assert output.size() == torch.Size([output_dim])
```

Note that each element of output scales as $\sqrt{\text{input_dim}}$: 18.919979095458984.

Large values can cause gradients to blow up and cause training to be unstable.

We want an initialization that is invariant to `input_dim`.

To do that, we simply rescale by $1/\sqrt{\text{input_dim}}$

```
w = nn.Parameter(torch.randn(input_dim, output_dim) / np.sqrt(input_dim))
```

```
output = x @ w # @inspect output
```

Now each element of output is constant: -1.5302726030349731.

Up to a constant, this is Xavier initialization. [\[paper\]](#)[\[stackexchange\]](#)

To be extra safe, we truncate the normal distribution to $[-3, 3]$ to avoid any chance of outliers.

```
w = nn.Parameter(nn.init.trunc_normal_(torch.empty(input_dim, output_dim), std=1 / np.sqrt(input_dim), a=-3, b=3))
```

```
def custom_model():
```

Let's build up a simple deep linear model using `nn.Parameter`.

```
D = 64 # Dimension
```

```
num_layers = 2
```

```
model = Cruncher(dim=D, num_layers=num_layers)
```

```
param_sizes = [
```

```
    (name, param.numel())
```

```
    for name, param in model.state_dict().items()
```

```
]
```

```
assert param_sizes == [
```

```
    ("layers.0.weight", D * D),
```

```
    ("layers.1.weight", D * D),
```

```
    ("final.weight", D),
```

```
]
```

```
num_parameters = get_num_parameters(model)
```

```
assert num_parameters == (D * D) + (D * D) + D
```

Remember to move the model to the GPU.

```
device = get_device()
```

```

629     model = model.to(device)
630
631     Run the model on some data.
632     B = 8 # Batch size
633     x = torch.randn(B, D, device=device)
634     y = model(x)
635     assert y.size() == torch.Size([B])
636
637
638 class Linear(nn.Module):
639     """Simple linear layer."""
640     def __init__(self, input_dim: int, output_dim: int):
641         super().__init__()
642         self.weight = nn.Parameter(torch.randn(input_dim, output_dim) / np.sqrt(input_dim))
643
644     def forward(self, x: torch.Tensor) -> torch.Tensor:
645         return x @ self.weight
646
647
648 class Cruncher(nn.Module):
649     def __init__(self, dim: int, num_layers: int):
650         super().__init__()
651         self.layers = nn.ModuleList([
652             Linear(dim, dim)
653             for i in range(num_layers)
654         ])
655         self.final = Linear(dim, 1)
656
657     def forward(self, x: torch.Tensor) -> torch.Tensor:
658         # Apply linear layers
659         B, D = x.size()
660         for layer in self.layers:
661             x = layer(x)
662
663         # Apply final head
664         x = self.final(x)
665         assert x.size() == torch.Size([B, 1])
666
667         # Remove the last dimension
668         x = x.squeeze(-1)
669         assert x.size() == torch.Size([B])
670
671         return x
672
673
674 def get_batch(data: np.array, batch_size: int, sequence_length: int, device: str) -> torch.Tensor:
675     Sample batch_size random positions into data.
676     start_indices = torch.randint(len(data) - sequence_length, (batch_size,))
677     assert start_indices.size() == torch.Size([batch_size])
678
679     Index into the data.
680     x = torch.tensor([data[start:start + sequence_length] for start in start_indices])
681     assert x.size() == torch.Size([batch_size, sequence_length])
682

```

Pinned memory

By default, CPU tensors are in paged memory. We can explicitly pin.

```

686 if torch.cuda.is_available():

```

```

687         x = x.pin_memory()
688
689     This allows us to copy x from CPU into GPU asynchronously.
690     x = x.to(device, non_blocking=True)
691
692     This allows us to do two things in parallel (not done here):
693     • Fetch the next batch of data into CPU
694     • Process x on the GPU.
695
696     \[article\]
697     \[article\]
698
699     return x
700
701
702 def note_about_randomness():
703     Randomness shows up in many places: parameter initialization, dropout, data ordering, etc.
704     For reproducibility, we recommend you always pass in a different random seed for each use of randomness.
705     Determinism is particularly useful when debugging, so you can hunt down the bug.
706
707     There are three places to set the random seed which you should do all at once just to be safe.
708
709     # Torch
710     seed = 0
711     torch.manual_seed(seed)
712
713     # NumPy
714     import numpy as np
715     np.random.seed(seed)
716
717     # Python
718     import random
719     random.seed(seed)
720
721
722 def data_loading():
723     In language modeling, data is a sequence of integers (output by the tokenizer).
724
725     It is convenient to serialize them as numpy arrays (done by the tokenizer).
726     orig_data = np.array([1, 2, 3, 4, 5, 6, 7, 8, 9, 10], dtype=np.int32)
727     orig_data.tofile("data.npy")
728
729     You can load them back as numpy arrays.
730     Don't want to load the entire data into memory at once (LLaMA data is 2.8TB).
731     Use memmap to lazily load only the accessed parts into memory.
732     data = np.memmap("data.npy", dtype=np.int32)
733     assert np.array_equal(data, orig_data)
734
735     A data loader generates a batch of sequences for training.
736     B = 2 # Batch size
737     L = 4 # Length of sequence
738     x = get_batch(data, batch_size=B, sequence_length=L, device=get_device())
739     assert x.size() == torch.Size([B, L])
740
741
742 class SGD(torch.optim.Optimizer):
743     def __init__(self, params: Iterable[nn.Parameter], lr: float = 0.01):
744         super(SGD, self).__init__(params, dict(lr=lr))

```

```

745
746     def step(self):
747         for group in self.param_groups:
748             lr = group["lr"]
749             for p in group["params"]:
750                 grad = p.grad.data
751                 p.data -= lr * grad
752
753
754 class AdaGrad(torch.optim.Optimizer):
755     def __init__(self, params: Iterable[nn.Parameter], lr: float = 0.01):
756         super(AdaGrad, self).__init__(params, dict(lr=lr))
757
758     def step(self):
759         for group in self.param_groups:
760             lr = group["lr"]
761             for p in group["params"]:
762                 # Optimizer state
763                 state = self.state[p]
764                 grad = p.grad.data
765
766                 # Get squared gradients  $g_2 = \sum_{i<t} g_i^2$ 
767                 g2 = state.get("g2", torch.zeros_like(grad))
768
769                 # Update optimizer state
770                 g2 += torch.square(grad)
771                 state["g2"] = g2
772
773                 # Update parameters
774                 p.data -= lr * grad / torch.sqrt(g2 + 1e-5)
775
776
777 def optimizer():
778     Recall our deep linear model.
779     B = 2
780     D = 4
781     num_layers = 2
782     model = Cruncher(dim=D, num_layers=num_layers).to(get_device())
783
784     Let's define the AdaGrad optimizer
785     • momentum = SGD + exponential averaging of grad
786     • AdaGrad = SGD + averaging by  $\text{grad}^2$ 
787     • RMSProp = AdaGrad + exponentially averaging of  $\text{grad}^2$ 
788     • Adam = RMSProp + momentum
789
790     AdaGrad: https://www.jmlr.org/papers/volume12/duchi11a/duchi11a.pdf
791     optimizer = AdaGrad(model.parameters(), lr=0.01)
792     state = model.state_dict() # @inspect state
793
794     Compute gradients
795     x = torch.randn(B, D, device=get_device())
796     y = torch.tensor([4., 5.], device=get_device())
797     pred_y = model(x)
798     loss = F.mse_loss(input=pred_y, target=y)
799     loss.backward()
800
801     Take a step
802     optimizer.step()

```

```
803 state = model.state_dict() # @inspect state
```

```
804
```

```
805 Free up the memory (optional)
```

```
806 optimizer.zero_grad(set_to_none=True)
```

```
807
```

808 Memory

```
809
```

```
810 # Parameters
```

```
811 num_parameters = (D * D * num_layers) + D # @inspect num_parameters
```

```
812 assert num_parameters == get_num_parameters(model)
```

```
813
```

```
814 # Activations
```

```
815 num_activations = B * D * num_layers # @inspect num_activations
```

```
816
```

```
817 # Gradients
```

```
818 num_gradients = num_parameters # @inspect num_gradients
```

```
819
```

```
820 # Optimizer states
```

```
821 num_optimizer_states = num_parameters # @inspect num_optimizer_states
```

```
822
```

```
823 # Putting it all together, assuming float32
```

```
824 total_memory = 4 * (num_parameters + num_activations + num_gradients + num_optimizer_states) # @inspect total_memory
```

```
825
```

826 Compute (for one step)

```
827 flops = 6 * B * num_parameters # @inspect flops
```

```
828
```

829 Transformers

```
830
```

```
831 The accounting for a Transformer is more complicated, but the same idea.
```

```
832 Assignment 1 will ask you to do that.
```

```
833
```

```
834 Blog post describing memory usage for Transformer training \[article\]
```

```
835 Blog post describing FLOPs for a Transformer: \[article\]
```

```
836
```

```
837
```

```
838 def train_loop():
```

```
839     Generate data from linear function with weights (0, 1, 2, ..., D-1).
```

```
840     D = 16
```

```
841     true_w = torch.arange(D, dtype=torch.float32, device=get_device())
```

```
842     def get_batch(B: int) -> tuple[torch.Tensor, torch.Tensor]:
```

```
843         x = torch.randn(B, D).to(get_device())
```

```
844         true_y = x @ true_w
```

```
845         return (x, true_y)
```

```
846
```

```
847     Let's do a basic run
```

```
848     train("simple", get_batch, D=D, num_layers=0, B=4, num_train_steps=10, lr=0.01)
```

```
849
```

```
850     Do some hyperparameter tuning
```

```
851     train("simple", get_batch, D=D, num_layers=0, B=4, num_train_steps=10, lr=0.1)
```

```
852
```

```
853
```

```
854 def train(name: str, get_batch,
```

```
855         D: int, num_layers: int,
```

```
856         B: int, num_train_steps: int, lr: float):
```

```
857     model = Cruncher(dim=D, num_layers=0).to(get_device())
```

```
858     optimizer = SGD(model.parameters(), lr=0.01)
```

```
859
```



```

860     for t in range(num_train_steps):
861         # Get data
862         x, y = get_batch(B=B)
863
864         # Forward (compute loss)
865         pred_y = model(x)
866         loss = F.mse_loss(pred_y, y)
867
868         # Backward (compute gradients)
869         loss.backward()
870
871         # Update parameters
872         optimizer.step()
873         optimizer.zero_grad(set_to_none=True)
874
875
876 def checkpointing():
877     Training language models take a long time and certainly will crash.
878     You don't want to lose all your progress.
879
880     During training, it is useful to periodically save your model and optimizer state to disk.
881
882     model = Cruncher(dim=64, num_layers=3).to(get_device())
883     optimizer = AdaGrad(model.parameters(), lr=0.01)
884
885     Save the checkpoint:
886     checkpoint = {
887         "model": model.state_dict(),
888         "optimizer": optimizer.state_dict(),
889     }
890     torch.save(checkpoint, "model_checkpoint.pt")
891
892     Load the checkpoint:
893     loaded_checkpoint = torch.load("model_checkpoint.pt")
894
895
896 def mixed_precision_training():
897     Choice of data type (float32, bfloat16, fp8) have tradeoffs.
898     • Higher precision: more accurate/stable, more memory, more compute
899     • Lower precision: less accurate/stable, less memory, less compute
900
901     How can we get the best of both worlds?
902
903     Solution: use float32 by default, but use {bfloat16, fp8} when possible.
904
905     A concrete plan:
906     • Use {bfloat16, fp8} for the forward pass (activations).
907     • Use float32 for the rest (parameters, gradients).
908
909     • Mixed precision training [Micikevicius+ 2017]
910
911     Pytorch has an automatic mixed precision (AMP) library.
912     https://pytorch.org/docs/stable/amp.html
913     https://docs.nvidia.com/deeplearning/performance/mixed-precision-training/
914
915     NVIDIA's Transformer Engine supports FP8 for linear layers
916     Use FP8 pervasively throughout training [Peng+ 2023]
917

```

```

918
919 #####
920
921 def get_memory_usage(x: torch.Tensor):
922     return x.numel() * x.element_size()
923
924
925 def get_promised_flop_per_sec(device: str, dtype: torch.dtype) -> float:
926     """Return the peak FLOP/s for `device` operating on `dtype`."""
927     if not torch.cuda.is_available():
928         No CUDA device available, so can't get FLOP/s.
929         return 1
930     properties = torch.cuda.get_device_properties(device)
931
932     if "A100" in properties.name:
933         # https://www.nvidia.com/content/dam/en-zz/Solutions/Data-Center/a100/pdf/nvidia-a100-datasheet-us-nvidia-
1758950-r4-web.pdf")
934         if dtype == torch.float32:
935             return 19.5e12
936         if dtype in (torch.bfloat16, torch.float16):
937             return 312e12
938         raise ValueError(f"Unknown dtype: {dtype}")
939
940     if "H100" in properties.name:
941         # https://resources.nvidia.com/en-us-tensor-core/nvidia-tensor-core-gpu-datasheet")
942         if dtype == torch.float32:
943             return 67.5e12
944         if dtype in (torch.bfloat16, torch.float16):
945             return 1979e12 / 2 # 1979 is for sparse, dense is half of that
946         raise ValueError(f"Unknown dtype: {dtype}")
947
948     raise ValueError(f"Unknown device: {device}")
949
950
951 def same_storage(x: torch.Tensor, y: torch.Tensor):
952     return x.untyped_storage().data_ptr() == y.untyped_storage().data_ptr()
953
954
955 def time_matmul(a: torch.Tensor, b: torch.Tensor) -> float:
956     """Return the number of seconds required to perform `a @ b`."""
957
958     # Wait until previous CUDA threads are done
959     if torch.cuda.is_available():
960         torch.cuda.synchronize()
961
962     def run():
963         # Perform the operation
964         a @ b
965
966         # Wait until CUDA threads are done
967         if torch.cuda.is_available():
968             torch.cuda.synchronize()
969
970     # Time the operation `num_trials` times
971     num_trials = 5
972     total_time = timeit.timeit(run, number=num_trials)
973
974     return total_time / num_trials

```

```
975
976
977 def get_num_parameters(model: nn.Module) -> int:
978     return sum(param.numel() for param in model.parameters())
979
980 def get_device(index: int = 0) -> torch.device:
981     """Try to use the GPU if possible, otherwise, use CPU."""
982     if torch.cuda.is_available():
983         return torch.device(f"cuda:{index}")
984     else:
985         return torch.device("cpu")
986
987 if __name__ == "__main__":
988     main()
```